

Tree Attention Forward Pass Latency: M4 Pro (MLX) vs A100 (CUDA)

